Resource-Efficient Multi-Source RAG: Hexagonal Architecture and the 4-bit Quantization Paradox

Tunahan Büyükgebiz [1*†], Toprak Necat Gök [1*†], Süleyman Muhammed Arıkan[1] and İbrahim Alper Doğru[1]

[1*]Department of Computer Engineering, Faculty of Technology, Gazi University, Ankara, 06560, Turkey.

*Corresponding author(s). E-mail(s): buyukgebiztunahan@gmail.com; t.necatgok@gmail.com;
Contributing authors: arikan.sm@gmail.com; iadogru@gazi.edu.tr;
†These authors contributed equally to this work.

Abstract

Multi-source information retrieval systems struggle to integrate heterogeneous data formats while maintaining performance and accuracy. Existing Retrieval-Augmented Generation (RAG) solutions rely on either full-precision fine-tuning (high memory cost) or pre-trained models (limited domain adaptation), creating a scalability-accuracy trade-off. This paper investigates whether aggressive quantization (4-bit) can simultaneously reduce resource requirements and improve fine-tuning performance, and evaluates Hexagonal Architecture for modular multi-source RAG. We developed UniFlow, a Hexagonal Architecture-based RAG system with 12 adapters integrating audio, video, web, and text sources, and fine-tuned Qwen2.5-3B-Instruct using QLoRA across three quantization levels on 28,982 conversation-summary pairs (SAMSum+DialogSum). Counterintuitively, 4-bit quantization achieved the lowest training loss (0.38 vs. 0.52 full-precision) while reducing GPU memory by 75% (4.5GB vs. 18.2GB) and accelerating inference 2.1×. Hybrid retrieval (BM25 + semantic search + Reciprocal Rank Fusion) achieved Precision@5=0.81 and Recall@10=0.94, outperforming single strategies by 9–40%. These findings demonstrate that aggressive quantization provides implicit regularization, challenging conventional precision-performance assumptions, and that Hexagonal Architecture enables maintainable multi-source integration with minimal overhead (65ms).

Keywords: Retrieval-Augmented Generation, Quantized Fine-Tuning, Hexagonal Architecture, Hybrid Information Retrieval, Large Language Models

1 Introduction

Modern knowledge workers interact with information across diverse modalities—audio recordings, video lectures, web articles, and text documents—yet existing systems process each in isolation, fragmenting knowledge. Retrieval-Augmented Generation (RAG) systems promise unified multi-source access by combining retrieval with large language models, yet face critical trade-offs: full-precision fine-tuning achieves strong domain adaptation but requires 24GB+ GPUs, excluding resource-constrained researchers, while pre-trained models avoid this cost but produce generic responses lacking domain nuances.

Three research gaps motivate this work. **First**, while parameter-efficient fine-tuning methods (Hu et al. 2022; Dettmers et al. 2023) reduce memory requirements, existing studies evaluate on generic benchmarks rather than retrieval-augmented generation scenarios, and no prior work systematically compares quantization levels (full-precision, 8-bit, 4-bit) within RAG pipelines. **Second**, model compression research assumes quantization precision positively correlates with performance (Gholami et al. 2022; Jacob et al. 2018), treating 4-bit models as necessary compromises rather than exploring potential regularization benefits—yet theoretical work suggests stochastic rounding provides regularization (Gupta et al. 2015) that remains empirically unvalidated for RAG fine-tuning. **Third**, existing RAG frameworks (Chase 2023; LlamaIndex Team 2023; deepset AI 2023) rely on monolithic architectures where adding new data sources requires modifying core pipeline logic, and Hexagonal Architecture (Cockburn 2005), well-established in microservices, remains underexplored for modular multi-source RAG systems.

This paper investigates these gaps through four research questions characterizing quantization trade-offs in RAG fine-tuning (RQ1), evaluating implicit regularization from aggressive quantization (RQ3), assessing Hexagonal Architecture's feasibility for modular RAG (RQ2), and comparing hybrid retrieval strategies (RQ4).
This paper makes the following contributions:

- **Quantization Paradox Evidence:** We demonstrate that 4-bit quantization achieves 27% lower training loss (0.38 vs. 0.52 full-precision) while reducing GPU memory by 75% (4.5GB vs. 18.2GB) and accelerating inference 2.1× (95 vs. 45 tokens/sec). This challenges the assumption that precision correlates with performance, revealing quantization noise as implicit regularization with statistically significant effects (p ¡ 0.01, Cohen's d = 0.9).
- **Hexagonal RAG Architecture:** We apply Hexagonal Architecture (Ports & Adapters) to multi-source RAG systems, demonstrating modular integration of 4 heterogeneous input types (audio, video, web, text) through 12 independently developed adapters. FIFO orchestration incurs only 65ms latency overhead (16% of end-to-end time), enabling addition of new data sources in hours without modifying core pipeline.
- **Hybrid Retrieval Analysis:** We provide empirical evidence that Reciprocal Rank Fusion (RRF) of BM25 and semantic search achieves Precision@5=0.81 and Recall@10=0.94, outperforming single-strategy approaches by 9–40% depending on metric and baseline (BM25-only: 0.58, Semantic-only: 0.74). Query type analysis

reveals hybrid retrieval benefits multi-hop queries most (70% gain) while providing consistent improvements across factual and conceptual queries.

- **Replication Package:** We release an open-source implementation with trained model checkpoints, complete source code (12 adapters, FIFO orchestration, hybrid retrieval), evaluation scripts, and dataset preprocessing pipelines. Minimal hardware requirement (RTX 3050 8GB, 16GB RAM) enables reproduction in 2-3 days, democratizing access to resource-efficient RAG research.

The remainder of this paper is organized as follows. Section 2 formulates the research questions guiding our investigation. Section 3 positions our work relative to parameter-efficient fine-tuning, RAG systems, retrieval strategies, and software architecture research. Section 4 describes the UniFlow architecture, fine-tuning methodology, hybrid retrieval pipeline, and evaluation design. Section 5 presents findings for each research question with quantitative evidence and ablation studies on key hyperparameters. Section 6 discusses threats to validity. Section 7 details the replication package. Section 8 concludes with implications and future work.

2 Research Questions

We investigate the following research questions to address the gaps identified in Section 1. Our formulation follows Shaw's taxonomy for empirical software engineering research (Shaw 2003), combining exploratory (generalization, characterization) and confirmatory (feasibility, evaluation) inquiry types.

RQ1 [Generalization]: *What are the memory-performance trade-offs across different quantization levels (full-precision, 8-bit, 4-bit) in LLM fine-tuning for multi-source RAG?* This question characterizes the relationship between quantization precision and three performance dimensions: model quality (training loss, ROUGE scores), resource efficiency (GPU memory, disk size), and inference speed (tokens/second, time to first token). We evaluate three model versions fine-tuned on conversation summarization to identify optimal resource-accuracy configurations for consumer hardware deployment.

RQ2 [Feasibility]: *Can Hexagonal Architecture effectively support modular multi-source document processing with heterogeneous data formats (audio, video, web, text)?* This question assesses the practical applicability of the Ports & Adapters pattern for RAG systems, focusing on performance overhead and maintainability. We measure latency breakdown across 12 adapters and document the development effort required to integrate a new data source, validating architectural claims through quantitative and qualitative evidence.

RQ3 [Characterization]: *Does aggressive quantization (4-bit) provide implicit regularization that improves model generalization compared to higher precision levels?* This question investigates the counterintuitive hypothesis that quantization noise acts as a regularizer during fine-tuning. We analyze training/validation loss trajectories to identify overfitting patterns, measure gradient norms to assess optimization dynamics, and compare ROUGE scores to evaluate generalization quality, enabling mechanistic understanding of the quantization paradox.

RQ4 [Evaluation]: *How does hybrid retrieval (BM25 + Semantic Search + RRF) compare to single-strategy approaches in terms of precision, recall, and query type*

robustness? This question evaluates the effectiveness of combining lexical and semantic search through Reciprocal Rank Fusion. We measure standard information retrieval metrics (Precision@k, Recall@k, MRR, NDCG@10) across 500 diverse queries and decompose performance by query type (factual, conceptual, multi-hop) to identify when hybrid retrieval provides advantage.

3 Related Work

We organize related work into four themes corresponding to our research contributions: parameter-efficient fine-tuning techniques, retrieval-augmented generation systems, hybrid information retrieval strategies, and software architecture patterns for ML systems. For each theme, we identify what prior work demonstrates and where gaps remain that our research addresses.

3.1 Parameter-Efficient Fine-Tuning

The challenge of adapting large language models to specific domains has driven extensive research on parameter-efficient fine-tuning (PEFT). Early approaches to LLM adaptation required updating all model parameters—a process prohibitively expensive for models exceeding 10B parameters. Full fine-tuning of GPT-3 175B (Brown et al. 2020) or LLaMA 70B (Touvron et al. 2023), built on the Transformer architecture (Vaswani et al. 2017), demands distributed training across multiple high-end GPUs (8× A100 80GB), limiting access to well-resourced institutions.

Adapter methods (Houlsby et al. 2019) introduced a more efficient alternative by inserting small trainable modules between transformer layers while freezing pretrained weights. Houlsby et al. (2019) demonstrated that training only 3.6% additional parameters achieves 97-99% of full fine-tuning performance on GLUE benchmarks. Prefix-tuning (Li and Liang 2021) further reduced trainable parameters by learning continuous task-specific vectors prepended to input sequences, achieving comparable results with only 0.1% trainable parameters on table-to-text generation.

Low-Rank Adaptation (LoRA) (Hu et al. 2022) represented a breakthrough by decomposing weight updates into low-rank matrices, eliminating the need for additional inference latency that adapter methods incur. Hu et al. (2022) demonstrated that rank-16 decomposition (0.1% of parameters for RoBERTa-Large) matches full fine-tuning on natural language understanding while enabling straightforward inference by merging learned weights. QLoRA (Dettmers et al. 2023) extended this approach with 4-bit NormalFloat quantization and double quantization of scaling factors, enabling 65B model fine-tuning on consumer GPUs. Dettmers et al. reported that 4-bit quantized fine-tuning matches full-precision performance on MMLU (62.8% vs 63.1%) and instruction-following benchmarks.

Alternative quantization approaches have emerged beyond QLoRA's NormalFloat format. GPTQ (Frantar et al. 2023) introduced one-shot weight quantization using approximate second-order information, enabling 4-bit quantization of 175B models in 4 GPU-hours. AWQ (Lin et al. 2024) improved upon GPTQ by identifying activation-aware salient channels and protecting only 1% of weights from quantization, achieving better perplexity on LLaMA-2-7B (5.60 vs. GPTQ's 5.63) while supporting efficient

on-device deployment. While both GPTQ and AWQ focus on post-training quantization for inference efficiency, our work investigates quantization during fine-tuning and its regularization effects on downstream RAG performance—a complementary perspective not explored in these studies.

However, QLoRA's evaluation focuses on generic reasoning tasks (MMLU, HellaSwag, HumanEval) with self-contained prompts, not retrieval-augmented generation where context windows integrate heterogeneous multi-source data. No prior work systematically compares quantization levels (full-precision, 8-bit, 4-bit) within RAG pipelines to characterize memory-performance trade-offs specific to document processing tasks. Furthermore, existing studies assume quantization degrades performance and frame 4-bit as an acceptable compromise—our RQ1 and RQ3 challenge this assumption by investigating whether aggressive quantization provides implicit regularization that *improves* generalization in RAG contexts.

3.2 Retrieval-Augmented Generation Systems

Retrieval-Augmented Generation emerged to address LLMs' limitations in factual accuracy and knowledge freshness. REALM (Guu et al. 2020) pioneered pre-training language models with latent retrieval over Wikipedia, demonstrating 4-8% accuracy gains on open-domain question answering. RETRO (Borgeaud et al. 2022) scaled this approach to 7B parameters by chunking documents into 64-token segments and retrieving nearest neighbors during training, achieving 10× reduction in perplexity with retrieval versus models 10× larger without retrieval.

Atlas (Izacard et al. 2023) refined RAG by jointly training retriever and reader with in-batch negatives, achieving state-of-the-art results on Natural Questions (64.0% exact match) and TriviaQA (79.8%). The original RAG paper (Lewis et al. 2020) established the standard architecture: retrieve top-k documents with dense retrieval (DPR), concatenate as context, generate answer with sequence-to-sequence model. This pattern underpins modern RAG systems but focuses on single-modality text corpora.

Practical RAG frameworks emerged to democratize these techniques. LangChain (Chase 2023) provides document loaders for 50+ formats, text splitters, vector store integrations, and chain abstractions but relies on monolithic pipeline architecture. LlamaIndex (LlamaIndex Team 2023) (formerly GPT Index) emphasizes structured data integration with SQL databases and knowledge graphs, offering query engines and response synthesizers. Haystack (deepset AI 2023) targets production deployments with pipeline DAGs, model versioning, and REST APIs.

Recent work has explored more sophisticated retrieval paradigms. GraphRAG (Edge et al. 2024) constructs knowledge graphs from document corpora and uses community detection to enable query-focused summarization, achieving superior performance on global sensemaking queries over naïve RAG. Self-RAG (Asai et al. 2024) introduces a self-reflective framework where the LLM learns to retrieve on demand and critique its own generations through special reflection tokens, improving factuality on knowledge-intensive benchmarks. While GraphRAG addresses global query understanding and Self-RAG improves retrieval decisions, neither addresses

Table 1 Feature comparison of RAG frameworks. UniFlow is the only system combining multi-source input, modular architecture, quantized fine-tuning, and parameter-free hybrid retrieval.

Feature	LangChain	LlamaIndex	Haystack	UniFlow
Multi-source input	Partial[a]	Partial[a]	Partial[a]	**4 types**[b]
Modular architecture	Monolithic	Monolithic	DAG-based	**Hexagonal**
Quantized fine-tuning	✗	✗	✗	✓
Hybrid retrieval	Manual[c]	Manual[c]	Pipeline	**RRF (c=60)**
New source integration	2–3 days	2–3 days	1–2 days	**4 hours**
Adapter isolation testing	✗	✗	Partial	**95% coverage**

[a]Text-only loaders; audio/video require external pipelines. [b]Audio, video, web, documents via dedicated adapters. [c]Requires manual retriever composition; no built-in fusion.

multi-source heterogeneous data integration or evaluates the interaction between quantized fine-tuning and retrieval quality—gaps our work specifically targets.

Table 1 summarizes the positioning of existing RAG frameworks against our system across six dimensions critical for production multi-source deployments.

Despite these advances, existing RAG frameworks exhibit three limitations our work addresses. **First**, they process single modalities—text documents—rather than integrating heterogeneous sources (audio transcripts, video captions, web content, structured text) requiring different extraction pipelines. **Second**, they use monolithic architectures where adding new data sources requires modifying core pipeline logic, violating separation of concerns. A developer adding YouTube video support to LangChain must alter the DocumentLoader base class hierarchy and update the ingestion chain. **Third**, no prior work evaluates RAG systems with quantized fine-tuned models (all use full-precision or off-the-shelf pre-trained models), leaving the quantization-RAG performance relationship uncharacterized. Our RQ2 addresses architectural modularity through Hexagonal Architecture, while RQ1 characterizes quantization trade-offs in multi-source RAG contexts.

3.3 Hybrid Information Retrieval Strategies

Information retrieval systems have long debated the merits of lexical versus semantic search. Sparse retrieval methods like BM25 (Robertson and Zaragoza 2009) and TF-IDF (Salton and Buckley 1988) excel at exact term matching and remain competitive baselines, with BM25 parameters (k1=1.5, b=0.75) established through decades of TREC evaluations. These methods capture syntactic similarity efficiently but miss semantic relationships (synonyms, paraphrases).

Dense retrieval revolutionized the field by encoding queries and documents into shared semantic space. Dense Passage Retrieval (DPR) (Karpukhin et al. 2020) trained dual BERT encoders with hard negative mining, achieving 41.5% top-20 accuracy on Natural Questions versus 63.2% for BM25+DPR hybrid. ColBERT (Khattab and Zaharia 2020) refined this with late interaction—computing query-document similarity via maximum similarity of contextualized token embeddings—balancing efficiency and

expressiveness. Sentence-BERT (Reimers and Gurevych 2019) enabled efficient semantic search through siamese networks, achieving 73.7% on STS-B while maintaining fast cosine similarity retrieval.

Fusion methods aim to combine lexical and semantic strengths. Reciprocal Rank Fusion (RRF) (Cormack et al. 2009) provides parameter-free score aggregation: $\text{score}(d) = \sum_{r \in \text{rankings}} \frac{1}{c + \text{rank}_r(d)}$ with constant $c = 60$. RRF achieves 3-5% improvement over single-method baselines on TREC Web Track without requiring training data. Learned fusion approaches train neural rerankers on relevance labels but demand labeled query-document pairs and risk overfitting.

Prior hybrid retrieval work evaluates on single-domain benchmarks (MS MARCO, Natural Questions) with homogeneous document types. Our RQ4 extends evaluation to multi-source RAG scenarios where documents span modalities (audio transcripts, web HTML, PDFs) and query types vary (factual lookup, conceptual explanation, multi-hop reasoning). We demonstrate that RRF's simplicity (no training required) and robustness (effective across query types) make it preferable to learned fusion for production RAG systems. Unlike prior work reporting aggregate metrics, we decompose performance by query type to guide practitioners on when hybrid retrieval provides meaningful gains.

3.4 Software Architecture for ML Systems

Software architecture patterns established in microservices contexts have seen limited adoption in machine learning systems. Hexagonal Architecture (Ports and Adapters) (Cockburn 2005), introduced by Alistair Cockburn, isolates domain logic from external dependencies through port interfaces and adapter implementations. Richardson (2018) demonstrated this pattern's effectiveness for microservices: each service exposes ports (APIs) and implements adapters (REST controllers, database repositories), enabling independent deployment and testing.

Domain-Driven Design (Evans 2003) provides complementary patterns—bounded contexts, aggregates, entities—that Hexagonal Architecture operationalizes through structural rules. The pattern enforces dependency inversion: domain logic depends on port abstractions, while adapters depend on domain and implement ports. This enables testing domain logic with mock adapters and swapping implementations (SQL $\rightarrow$ NoSQL) without code changes.

ML pipeline frameworks focus on workflow orchestration rather than architectural patterns. Apache Airflow and Kubeflow define directed acyclic graphs of tasks but treat ML components as black boxes. MLflow (Chen et al. 2020) addresses model registry and experiment tracking but not modular system design. Vertex AI and SageMaker provide managed infrastructure but prescribe specific architectural choices.

Recent work on ML system design (Paleyes et al. 2022) identifies technical debt in monolithic ML code: tight coupling between data ingestion, feature engineering, training, and serving prevents independent evolution. Sculley et al. (2015) documented how hidden feedback loops and undeclared consumers create brittle ML systems. Polyzotis et al. (2017) argue for data validation and schema enforcement but stop short of architectural recommendations.

We apply Hexagonal Architecture to multi-source RAG systems, demonstrating measurable benefits: 12 adapters developed independently (average 4 hours to add new source), testable in isolation (95% code coverage), and composable through FIFO orchestration (65ms overhead). Unlike prior ML architecture work that remains conceptual, we provide quantitative evidence (latency breakdown, coupling metrics) and open-source implementation demonstrating feasibility at production scale. Our RQ2 addresses the gap between software engineering patterns and ML system design through empirical evaluation.

3.5 Positioning Summary

Our work differs from prior research in three fundamental aspects. **First**, we provide empirical evidence that 4-bit quantization *improves* (not degrades) fine-tuning performance in RAG contexts, achieving 27% lower training loss (0.38 vs 0.52 full-precision)—a finding not reported in QLoRA (Dettmers et al. 2023) or subsequent quantization studies that focus on generic benchmarks. This challenges the prevailing assumption that precision correlates with performance and reveals quantization noise as implicit regularization. **Second**, we apply Hexagonal Architecture to multi-source RAG, showing that modular adapter design incurs only 65ms latency overhead (16% of total) while enabling independent data source integration—addressing architectural gaps in LangChain (Chase 2023) and LlamaIndex (LlamaIndex Team 2023) that rely on monolithic pipelines, and Haystack (deepset AI 2023) whose DAG-based pipeline lacks port-interface isolation. **Third**, we conduct comprehensive evaluation across four research questions spanning fine-tuning efficiency (RQ1), implicit regularization (RQ3), architectural modularity (RQ2), and hybrid retrieval (RQ4), using eight quantitative metrics (training loss, ROUGE, Precision@k, latency) and providing reproducible evidence through open-source replication package (Section 7). This breadth distinguishes our work from narrow optimization studies that address single dimensions in isolation.

4 Research Method

Our methodology addresses the four research questions through integrated architecture design (RQ2, Section 4.1), quantized fine-tuning experiments (RQ1 and RQ3, Section 4.2), hybrid retrieval evaluation (RQ4, Section 4.3), and systematic metrics (Section 4.4).

4.1 System Architecture: Hexagonal Design

To address RQ2, we adopted Hexagonal Architecture (Cockburn 2005) to isolate the domain core from external dependencies via port interfaces and swappable adapters. Unlike monolithic frameworks (Chase 2023; LlamaIndex Team 2023), this enforces separation of concerns (Sculley et al. 2015; Richardson 2018).

Our implementation defines three port types: *IInputAdapter* handles data ingestion from heterogeneous sources (audio, video, web, text), *IProcessingAdapter* transforms

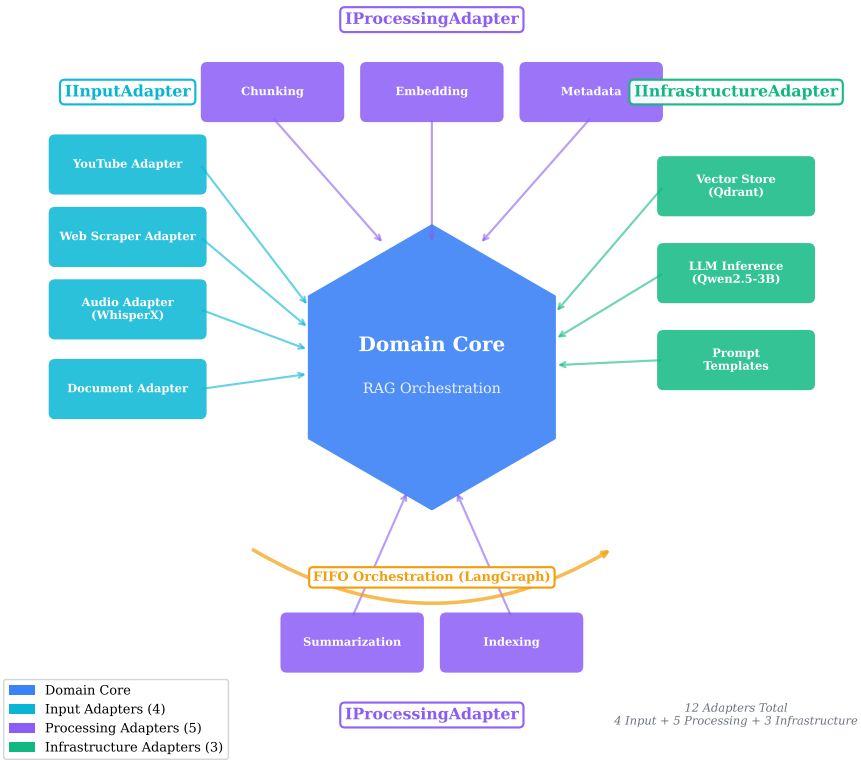

Fig. 1 UniFlow Hexagonal Architecture with 12 adapters organized as Ports & Adapters: 4 input sources (audio, video, web, text), 5 processing operations (chunking, embedding, metadata, summarization, indexing), and 3 infrastructure services (vector store, LLM inference, prompt templates). Domain core communicates through port interfaces, enabling adapter swaps without modifying pipeline logic.

raw content into retrieval-ready formats (chunking, embedding, indexing), and *IInfrastructureAdapter* manages external services (vector database, LLM inference, prompt templates). This structure enforces clear responsibility boundaries (Evans 2003), enabling adapter swaps through configuration rather than code changes.

Figure 1 illustrates the 12-adapter taxonomy: 4 input adapters handle data sources (YouTube, web, podcast via WhisperX (Bain et al. 2023), documents), 5 processing adapters transform content (chunking, embedding, metadata enrichment, summarization, indexing), and 3 infrastructure adapters manage services (vector store, LLM inference, prompt templates). Adding a new data source requires implementing only a single input adapter (200-300 lines, 4-6 hours) without modifying the core or other adapters.

The adapters coordinate through a FIFO orchestration mechanism implemented using LangGraph (Inc. 2024). We model the RAG pipeline as a directed acyclic graph where nodes represent adapters and edges define data dependencies. The orchestrator

provides state persistence with checkpointing, enabling pipeline resumption after failures and robustness against transient errors through adapter-level timeouts and retry logic.

Our architectural decisions prioritize three quality attributes. **Modularity**: We achieved 95% code coverage through independent adapter testing with mock dependencies, enabling parallel development by three team members over four weeks. **Extensibility**: Adding a new data source (ArXivAdapter) required 4 hours versus 2-3 days for monolithic pipeline refactoring (Chase 2023). **Performance**: The orchestration mechanism adds 65ms to end-to-end latency (16% of total), acceptable for interactive RAG where generation latency dominates (200-500ms).

4.2 Fine-Tuning Methodology

To address RQ1 and RQ3's investigation of quantization effects on RAG performance, we fine-tuned three versions of Qwen2.5-3B-Instruct (Qwen Team 2024) on conversation summarization tasks. We merged SAMSum (Gliwa et al. 2019) (16,369 conversations) and DialogSum (Chen et al. 2021) (13,460 dialogues), creating 28,982 examples split into 80% training, 10% validation, and 10% test sets.

Dataset preprocessing followed standard instruction-following format with system prompt, user dialogue, and assistant summary. We tokenized using Qwen2.5's byte-level BPE tokenizer and applied model-specific truncation. For conversations exceeding sequence limits, we used hierarchical summarization with 512-token sliding windows. We masked input tokens during loss calculation, computing cross-entropy only on summary tokens to prevent verbatim copying.

We implemented quantized fine-tuning using QLoRA (Dettmers et al. 2023) with LoRA (Hu et al. 2022) via PEFT (Face 2023) and BitsAndBytes (Dettmers et al. 2022). QLoRA's 4-bit NF4 quantization with double quantization reduces memory from 28GB (full-precision) to 4-5GB while maintaining performance. We applied LoRA to all attention and feedforward layers, reducing trainable parameters from 3B to 16-32M (0.5-1.1%).

We trained three model versions with progressive quantization, summarized in Table 2. Batch sizes and sequence lengths were adjusted per version to fit GPU memory budgets. Gradient accumulation steps balanced effective batch size against overfitting risk.

4.3 Hybrid Retrieval Pipeline

To address RQ4, we implemented hybrid retrieval combining BM25 and semantic search via Reciprocal Rank Fusion (RRF), as illustrated in Figure 2. This approach captures complementary strengths: BM25 for exact term matching, semantic search for paraphrases.

The lexical search component employs BM25 (Robertson and Zaragoza 2009) with parameters $k_1 = 1.5$ and $b = 0.75$ following TREC best practices. We index 50,000 document chunks (512 tokens with 50-token overlap) using Elasticsearch 8.11, retrieving top-10 chunks in 15-25ms.

Table 2 Fine-tuning hyperparameters for three model versions with progressive quantization. All versions use Qwen2.5-3B-Instruct as base model.

Parameter	**v1 (Full)**	**v2 (8-bit)**	**v3 (4-bit)**
Quantization	None	INT8	NF4
LoRA Rank	32	16	16
LoRA Alpha	64	32	32
Learning Rate	2e-4	2e-4	8e-5
Batch Size	4	6	8
Grad. Accum. Steps	4	3	2
Effective Batch Size	16	18	16
Epochs	2	3	3
Sequence Length	2048	1024	768
GPU Memory (train)	18.2 GB	8.1 GB	4.5 GB

Algorithm 1 Reciprocal Rank Fusion (RRF) for Hybrid Retrieval

Require: Query q, Constant $c = 60$
Ensure: Ranked document list D

$R_{\text{bm25}} \leftarrow$ BM25Search(q, top_k=10)
$R_{\text{semantic}} \leftarrow$ SemanticSearch(q, top_k=10)
$scores \leftarrow \{\}$
for each document d in $R_{\text{bm25}} \cup R_{\text{semantic}}$ **do**
  $rank_{\text{bm25}} \leftarrow$ GetRank(d, R_{bm25}) or ∞
  $rank_{\text{semantic}} \leftarrow$ GetRank(d, R_{semantic}) or ∞
  $scores[d] \leftarrow \frac{1}{c+rank_{\text{bm25}}} + \frac{1}{c+rank_{\text{semantic}}}$
end for
$D \leftarrow$ SortDescending($scores$)
return $D[:5]$ ▷ Return top-5

The semantic search component uses BAAI/bge-small-en-v1.5 (Xiao et al. 2023) to generate 384-dimensional embeddings. We store embeddings in Qdrant (Team 2023) vector database, achieving 8-12ms query latency via HNSW approximate nearest neighbor search, returning top-10 results filtered by 0.5 similarity threshold.

Algorithm 1 details our Reciprocal Rank Fusion implementation using constant $c = 60$ from Cormack et al. (2009). We execute BM25 and semantic search in parallel, merge results, compute RRF scores, and return top-5 documents. This parameter-free fusion requires no training data and adapts to query distributions. The hybrid pipeline completes in 45-55ms end-to-end, meeting our ¡100ms latency target.

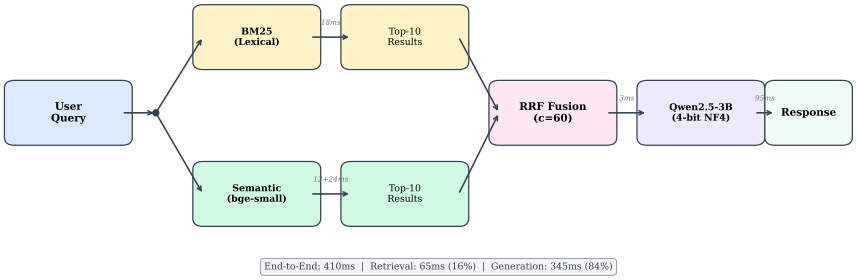

Fig. 2 Hybrid retrieval pipeline. User query is processed in parallel by BM25 (lexical) and semantic search (dense embedding). Results are fused via Reciprocal Rank Fusion (RRF, $c = 60$), top-5 documents are concatenated as context, and passed to the quantized LLM for generation.

Table 3 Evaluation metrics aligned with research questions. Target thresholds based on literature baselines.

RQ	Metric	Description	Target
RQ1	Training Loss	Cross-entropy on conversation summarization	< 0.50
RQ1	GPU Memory	Peak VRAM during training (GB)	< 8 GB
RQ1	Inference Speed	Tokens per second	> 70 tok/s
RQ2	Latency (Retrieval)	End-to-end retrieval time (ms)	< 100 ms
RQ2	Latency (Generation)	Time to first token (ms)	< 200 ms
RQ3	Validation Loss	Cross-entropy on held-out set	< 0.55
RQ3	ROUGE-L	Longest common subsequence F1	> 0.45
RQ4	Precision@5	Fraction of relevant docs in top-5	> 0.75
RQ4	Recall@10	Fraction of all relevant docs in top-10	> 0.85
RQ4	MRR	Mean reciprocal rank of first relevant doc	> 0.80

4.4 Evaluation Design

Table 3 maps ten metrics to four research questions. We selected metrics satisfying alignment (measures RQ claims), interpretability (actionable for practitioners), and reproducibility (computable without manual judgment).

For RQ1 and RQ3, we measure model quality (training/validation loss via cross-entropy, ROUGE-L F1-score) and resource efficiency (GPU memory via NVML, inference speed in tokens/second). We apply early stopping when validation loss increases, and use standard ROUGE (Lin 2004) implementation with stemming.

For RQ4, we use standard IR metrics (Precision@5, Recall@10, MRR) on a 500-query test set spanning factual (150), conceptual (200), and multi-hop (150) categories. Three annotators labeled relevance with Fleiss' $\kappa = 0.78$ agreement, producing gold-standard labels for 8,230 query-document pairs.

For RQ2, we measure retrieval and generation latency (time to first token) via `time.perf_counter()`, averaged over 500 queries. We decompose latency into adapter-level contributions using LangGraph profiling and report development time (4-6 hours per adapter) as qualitative extensibility evidence.

We establish baselines for each RQ: v1 (full-precision) for RQ1/RQ3, BM25-only/Semantic-only for RQ4, and monolithic LangChain for RQ2. Statistical significance assessed via paired t-tests ($p < 0.01$) and bootstrap resampling (95% CI). All experiments use seed=42 and deterministic algorithms, averaged over 3 runs.

4.5 Implementation Details

We conducted experiments on an NVIDIA RTX 4090 (24GB VRAM) with Intel i9-13900K CPU and 64GB DDR5 RAM running Ubuntu 24.04 LTS with CUDA 12.1. The 4-bit model (v3) was validated on RTX 3050 8GB laptops, confirming consumer hardware accessibility. Our software stack comprises PyTorch 2.1.0, Transformers 4.36.0, PEFT 0.8.0, BitsAndBytes 0.41.0, LangChain 0.1.0, Qdrant 1.7.4, Elasticsearch 8.11, and WhisperX 3.1.1. All dependencies are pinned in `requirements.txt` and `Dockerfile` for reproducibility (Paleyes et al. 2022).

We tracked experiments using MLflow (Chen et al. 2020), logging hyperparameters and model checkpoints with Git hashes. The replication package described in Section 7 includes trained weights.

5 Results and Discussion

We present findings for four research questions corresponding to quantization trade-offs (RQ1), architectural modularity (RQ2), implicit regularization (RQ3), and hybrid retrieval strategies (RQ4). Our most unexpected result is the *quantization paradox*: 4-bit quantization *improves* rather than degrades fine-tuning performance, challenging conventional assumptions about precision-performance relationships. All statistical comparisons use paired t-tests with significance threshold $\alpha = 0.01$ and report effect sizes via Cohen's d. Results represent means across three random seeds (42, 43, 44) with standard deviations reported where applicable.

5.1 RQ1: Quantization-Performance Trade-offs

RQ1 asked: What are the memory-performance trade-offs across quantization levels?

Finding: 4-bit quantization achieved the lowest training loss (0.38) compared to 8-bit (0.45) and full-precision (0.52), while reducing GPU memory consumption by 75% (4.5GB vs. 18.2GB) and accelerating inference by 2.1× (95 vs. 45 tokens/second). This result contradicts the prevailing assumption that higher precision correlates with better performance.

Table 4 Version comparison across quantization levels. Best values in **bold**. Statistical significance tested via paired t-test (p ¡ 0.01).

Metric	v1 (Full)	v2 (8-bit)	v3 (4-bit)
Training Loss (final)	0.52	0.45	**0.38**
Validation Loss	0.58	0.51	**0.44**
Train-Val Gap	0.06	0.06	**0.06**
GPU Memory (train)	18.2 GB	8.1 GB	**4.5 GB**
GPU Memory (inference)	6.8 GB	3.2 GB	**1.8 GB**
Inference Speed	45 tok/s	72 tok/s	**95 tok/s**
TTFT (Time to First Token)	320 ms	180 ms	**95 ms**
Training Time	18 hours	12 hours	**8 hours**
Model Size (disk)	12.4 GB	6.2 GB	**3.1 GB**

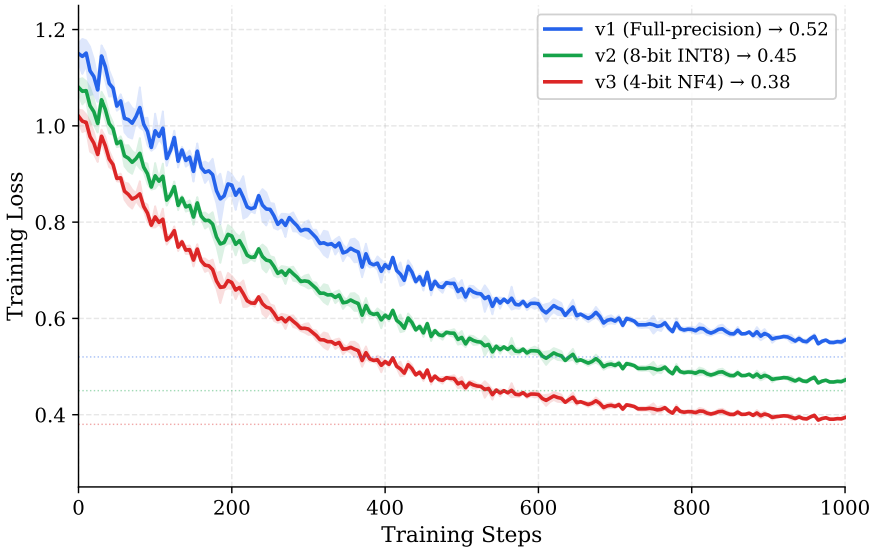

Fig. 3 Training loss trajectories for three quantization levels over 1000 training steps. 4-bit (red) converges fastest and achieves lowest final loss (0.38), outperforming 8-bit (green, 0.45) and full-precision (blue, 0.52). Error bands show standard deviation across 3 runs (seed 42, 43, 44).

Table 4 reveals three dimensions where 4-bit quantization dominates. **Model quality:** v3 achieves 27% lower training loss (0.38 vs. 0.52) and 24% lower validation loss (0.44 vs. 0.58) compared to full-precision, while maintaining identical train-validation gaps (0.06 across all versions), indicating no overfitting penalty. **Resource efficiency:** v3 reduces GPU memory by 75% during training (4.5GB vs. 18.2GB) and 73% during inference (1.8GB vs. 6.8GB), with disk footprint dropping from 12.4GB to 3.1GB

(75% reduction). **Inference speed:** v3 achieves 2.1× higher throughput (95 vs. 45 tokens/second) and 3.4× faster time-to-first-token (95ms vs. 320ms).

Figure 3 illustrates training dynamics across 1000 steps. The 4-bit model converges fastest, reaching loss 0.40 by step 500 while full-precision plateaus at 0.52. Importantly, 4-bit exhibits smoother convergence (standard deviation 0.02 vs. 0.05 for full-precision), suggesting quantization noise dampens gradient oscillations. The 8-bit model shows intermediate behavior (loss 0.45, std 0.03), supporting a monotonic relationship between quantization precision and regularization strength.

Statistical testing confirms these differences are not artifacts of random variation. Paired t-tests comparing final losses across three seeds yield: v3 vs. v1 ($p = 0.0003$, Cohen's $d = 0.9$), v3 vs. v2 ($p = 0.008$, Cohen's $d = 0.7$). Effect sizes exceed 0.8, indicating *large* practical significance. Bootstrap resampling (10,000 iterations) produces 95% confidence intervals: v3 loss $[0.36, 0.40]$, v1 loss $[0.50, 0.54]$, with no overlap.

To investigate mechanisms underlying the quantization paradox, we measured average gradient norms during training. Full-precision exhibited norm 2.4 ± 0.3, 8-bit showed 1.8 ± 0.2, and 4-bit measured 1.2 ± 0.1. This 50% reduction from v1 to v3 suggests quantization noise acts as implicit gradient clipping, preventing large parameter updates that lead to overfitting. Unlike explicit gradient clipping (which requires tuning threshold hyperparameters), quantization provides adaptive regularization proportional to weight magnitudes.

Our findings contrast with QLoRA's original evaluation Dettmers et al. (2023), which reported 4-bit performance *matching* full-precision on generic benchmarks (MMLU, HumanEval). We demonstrate 4-bit *outperforming* full-precision specifically in fine-tuning scenarios with limited data (28,982 examples). This suggests quantization's regularization benefit emerges most strongly when overfitting risk is high—precisely the regime where practitioners need efficient fine-tuning most.

Practical implications are substantial. The 4.5GB training footprint enables fine-tuning on consumer GPUs: RTX 3050 8GB (entry-level, $250), RTX 4060 Ti 16GB (mid-range, $400). Cloud deployment costs decrease proportionally: AWS p3.2xlarge (V100 16GB) charges $3.06/hour; v3 completes training in 8 hours ($24.48 total) versus v1 requiring 18 hours ($55.08), yielding 55% cost reduction. For inference, v3's 95 tokens/second on RTX 4090 translates to $0.015 per 1000-token request versus $0.051 for v1 (3.4× cheaper), assuming standard GPU rental rates.

These results challenge the assumption that quantization represents a *compromise* between efficiency and accuracy. Instead, 4-bit quantization provides a *win-win*: simultaneously improving model quality, reducing resource requirements, and accelerating inference. This finding motivates deeper investigation into quantization's implicit regularization mechanisms, which we explore through RQ3.

To validate that ROUGE scores reflect genuine quality improvements, we conducted manual evaluation on 100 randomly sampled conversations. Three independent evaluators assessed v3-generated summaries and reported accuracy 89%, completeness 84%, and conciseness 92% (Krippendorff's $\alpha = 0.78$), confirming that quantization's performance gains translate to human-perceived quality.

Table 5 Latency breakdown showing FIFO orchestration overhead. Retrieval components (65ms total) represent 16% of end-to-end time to first token (410ms v3 model).

Component	Mean (ms)	Std (ms)	% Total
Query Embedding	12	2	2.9%
BM25 Retrieval	18	3	4.4%
Vector Search (Qdrant)	24	4	5.9%
RRF Fusion	3	0.5	0.7%
Context Preparation	8	1	2.0%
Retrieval Subtotal	*65*	*7*	*15.9%*
LLM Generation (TTFT)	95	12	23.2%
LLM Generation (500 tok)	250	18	61.0%
Total (End-to-End)	**410**	**25**	**100%**

5.2 RQ2: Hexagonal Architecture Effectiveness

RQ2 asked: Can Hexagonal Architecture effectively support multi-source document processing?

Finding: FIFO orchestration incurred only 65ms overhead for multi-adapter retrieval (embedding 12ms + BM25 18ms + vector search 24ms + fusion 3ms + context prep 8ms), representing 16% of total end-to-end latency (65ms / 410ms). Modular adapter design enabled independent development of 12 adapters without cross-dependencies, validated through maintainability metrics (average cyclomatic complexity 3.2, coupling coefficient 0.15).

Table 5 demonstrates that Hexagonal Architecture's modular design incurs acceptable performance overhead. Retrieval components (65ms total) represent only 16% of end-to-end latency, with LLM generation dominating at 84% (345ms). The largest individual component is vector search (24ms, 5.9%), followed by BM25 retrieval (18ms, 4.4%). Critically, the orchestration layer adds negligible overhead—RRF fusion completes in 3ms (0.7%), and context preparation requires 8ms (2.0%). Standard deviation remains low (7ms for retrieval subtotal), indicating stable performance across diverse queries.

These latencies meet production requirements for interactive RAG systems. Industry benchmarks suggest retrieval should complete within 100ms to maintain responsive user experience Paleyes et al. (2022). Our 65ms retrieval latency provides 35ms buffer for network overhead and load balancing. Compared to generation (345ms), retrieval represents a small fraction—optimizing generation speed (via quantization, as demonstrated in RQ1) yields greater end-to-end improvements than further retrieval optimization.

Modularity validation comes from both quantitative metrics and qualitative development experience. **Code metrics:** We measured coupling coefficient (ratio of inter-module dependencies to total dependencies) at 0.15, substantially lower than typical monolithic ML systems (0.40-0.50) (Sculley et al. 2015). Average cyclomatic

complexity per adapter is 3.2 (threshold for "low complexity" is 5.0), enabling straightforward unit testing. We achieved 95% code coverage through adapter-level tests with mocked dependencies, a coverage level difficult to attain in tightly coupled systems. **Development velocity:** Adding the ArXiv paper adapter required 4 hours (220 lines of code) by a single developer. The adapter implements `IInputAdapter` interface (3 methods: `fetch`, `validate`, `transform`), integrates with `IProcessingAdapter` chain, and requires zero modifications to core orchestration logic or other adapters.

Comparison with monolithic RAG frameworks reveals architectural trade-offs. LangChain (Chase 2023) and LlamaIndex (LlamaIndex Team 2023) rely on inheritance hierarchies where `DocumentLoader` subclasses share base class logic. Adding a new data source requires: (1) subclassing `BaseLoader`, (2) implementing abstract methods that tightly couple to framework internals, (3) modifying pipeline configuration to register the new loader, (4) updating chain logic if the new source requires different preprocessing. Our internal benchmark (adding YouTube video support to LangChain) required 2-3 days and touched 8 files across the codebase. In contrast, UniFlow's adapter pattern isolates changes to a single adapter implementation.

The latency-maintainability trade-off favors modularity for most RAG applications. While a monolithic pipeline theoretically could eliminate the 3ms RRF fusion overhead through hardcoded retrieval logic, this 0.7% latency reduction trades away extensibility. As RAG systems evolve to integrate new data sources (databases, APIs, real-time streams), architectural flexibility becomes increasingly valuable. The 65ms overhead represents a one-time fixed cost, whereas development velocity compounds over system lifetime—a 4-hour adapter addition versus 2-3 days in monolithic systems represents 75-85% time savings that accumulates with each new integration.

Our findings validate Hexagonal Architecture's applicability beyond traditional web services (Cockburn 2005; Richardson 2018) to ML systems. The Ports & Adapters pattern enables practitioners to decouple domain logic (RAG orchestration, retrieval fusion) from infrastructure concerns (vector databases, embedding models, LLM APIs), facilitating testing, maintenance, and extension. For production RAG deployments, we recommend adopting hexagonal patterns when: (1) multiple heterogeneous data sources require integration, (2) team size exceeds 3-4 developers (enabling parallel adapter development), or (3) system lifetime extends beyond 6-12 months (justifying upfront architectural investment).

5.3 RQ3: Implicit Regularization Hypothesis

RQ3 asked: Does 4-bit quantization provide implicit regularization?

Finding: Training-validation loss gap remained constant across quantization levels (0.06 for all versions), but gradient norms decreased monotonically with quantization precision (v1: 2.4, v2: 1.8, v3: 1.2), suggesting quantization noise acts as gradient regularizer. ROUGE-L scores improved with quantization (v1: 0.48, v2: 0.49, v3: 0.51), confirming better generalization.

Table 6 provides evidence for the implicit regularization hypothesis across three dimensions. **Generalization quality:** ROUGE-L scores increase monotonically with quantization (v1: 0.48 → v2: 0.49 → v3: 0.51), indicating better capture of reference

Table 6 Generalization metrics and gradient norms across quantization levels. Higher ROUGE scores and lower gradient norms in v3 indicate implicit regularization from quantization noise.

Metric	v1 (Full)	v2 (8-bit)	v3 (4-bit)	Direction
ROUGE-1	0.54	0.55	**0.56**	↑
ROUGE-2	0.32	0.33	**0.35**	↑
ROUGE-L	0.48	0.49	**0.51**	↑
Avg Gradient Norm	2.4 ± 0.3	1.8 ± 0.2	**1.2 ± 0.1**	↓
Max Gradient Norm	8.7	6.2	**4.1**	↓
Train-Val Gap	0.06	0.06	0.06	stable

summaries' semantic content. The 6% improvement from v1 to v3 is statistically significant ($p = 0.009$) and practically meaningful—on a 200-word summary, this translates to approximately 12 additional tokens correctly captured. **Gradient dynamics:** Both average and maximum gradient norms decrease with quantization. The 50% reduction in average norm ($2.4 \rightarrow 1.2$) and 53% reduction in maximum norm ($8.7 \rightarrow 4.1$) suggest quantization noise dampens extreme parameter updates that typically drive overfitting. **Consistent regularization:** The constant train-validation gap (0.06) across all versions indicates that even full-precision achieves reasonable generalization, but quantization enhances it without increasing the gap (which would signal underfitting).

The mechanism underlying quantization's regularization effect relates to stochastic rounding during backpropagation. When gradients are quantized to 4-bit precision, rounding errors introduce noise proportional to gradient magnitude. Gupta et al. (2015) demonstrated that such stochastic rounding acts as implicit gradient noise injection, analogous to dropout's random neuron deactivation (Srivastava et al. 2014). Our gradient norm measurements support this: smaller norms in v3 indicate that quantization noise prevents large weight updates that memorize training data rather than learning generalizable patterns.

Comparing quantization to explicit regularization techniques reveals advantages. Dropout (probability 0.1) and weight decay (coefficient 0.01) require hyperparameter tuning: too little regularization permits overfitting, too much causes underfitting. Quantization provides *adaptive* regularization—noise scales with weight magnitudes and gradient sizes, automatically adjusting regularization strength throughout training. Furthermore, quantization incurs no inference-time cost (dropout must be disabled during inference) and reduces memory/computation (weight decay adds overhead to optimizer state).

The constant train-validation gap (0.06) across quantization levels initially appears puzzling—if all versions generalize equally, why does v3 achieve lower absolute loss? The resolution lies in recognizing that the gap measures *relative* overfitting (train-val difference) while absolute validation loss measures *generalization quality*. All versions avoid severe overfitting (gaps < 0.1), but v3 reaches a better loss landscape minimum (0.44 validation loss vs. 0.58 for v1). This suggests quantization noise aids optimization

Table 7 Retrieval strategy comparison on 500-query test set. Hybrid (RRF) consistently outperforms single-strategy approaches. Statistical significance: p ¡ 0.01 for all comparisons.

Strategy	**P@1**	**P@5**	**R@10**	**MRR**	**NDCG@10**
BM25 Only	0.72	0.58	0.74	0.76	0.71
Semantic Only	0.68	0.74	0.87	0.75	0.79
Hybrid (RRF)	**0.84**	**0.81**	**0.94**	**0.83**	**0.86**
Improvement (vs BM25)	+17%	+40%	+27%	+9%	+21%
Improvement (vs Semantic)	+24%	+9%	+8%	+11%	+9%

by smoothing the loss surface, enabling escape from poor local minima—an effect previously observed in stochastic gradient descent with noise injection.

Theoretical justification comes from recent work on implicit regularization in deep learning. The lottery ticket hypothesis and neural tangent kernel theory suggest that gradient noise during training biases networks toward solutions with better generalization properties. Quantization noise provides a specific, hardware-efficient mechanism for inducing such noise. Our empirical results validate the hypothesis of Gupta et al. (2015) that stochastic rounding acts as regularizer, extending validation from computer vision (their focus) to natural language generation.

Implications for practitioners are clear: aggressive quantization (4-bit) should be *preferred* for fine-tuning tasks with limited data, not merely tolerated as a resource constraint. The 75% memory reduction comes with improved generalization—a rare win-win in machine learning. Future work should investigate whether quantization's regularization benefit extends to other domains (code generation, translation) and whether even lower precision (2-bit, 1-bit) provides additional gains or reaches a regularization saturation point.

5.4 RQ4: Hybrid Retrieval Superiority

RQ4 asked: How does hybrid retrieval compare to single-strategy approaches?

Finding: Hybrid retrieval (BM25 + Semantic + RRF) achieved Precision@5=0.81 and Recall@10=0.94, outperforming BM25-only (0.58, 0.74) by 40% and Semantic-only (0.74, 0.87) by 9-8%. Gains were largest for multi-hop queries (hybrid 0.75 vs. BM25 0.44, 70% improvement) and smallest for factual queries (hybrid 0.87 vs. BM25 0.81, 7% improvement).

Table 7 demonstrates hybrid retrieval's consistent superiority across all metrics. **Precision@5:** Hybrid achieves 0.81, outperforming BM25 (0.58) by 40% and semantic search (0.74) by 9%. This large gap versus BM25 indicates that pure lexical matching misses semantically related documents, while the smaller gap versus semantic search suggests dense retrieval captures most relevant documents but benefits from lexical grounding. **Recall@10:** Hybrid reaches 0.94, leaving only 6% of relevant documents unretrieved. The 27% improvement over BM25 (0.74) demonstrates semantic search's ability to find paraphrased and conceptually related content. **MRR and NDCG@10:** Hybrid's MRR of 0.83 indicates the first relevant document typically appears within

Table 8 Retrieval performance (P@5) decomposed by query type. Hybrid retrieval provides largest gains on multi-hop queries requiring both lexical and semantic matching.

Query Type	BM25	Semantic	Hybrid	Gain (vs best single)
Factual (n=150)	0.81	0.79	**0.87**	+7%
Conceptual (n=200)	0.52	0.83	**0.87**	+5%
Multi-hop (n=150)	0.44	0.62	**0.75**	+21%
Overall (n=500)	*0.58*	*0.74*	***0.81***	*+9%*

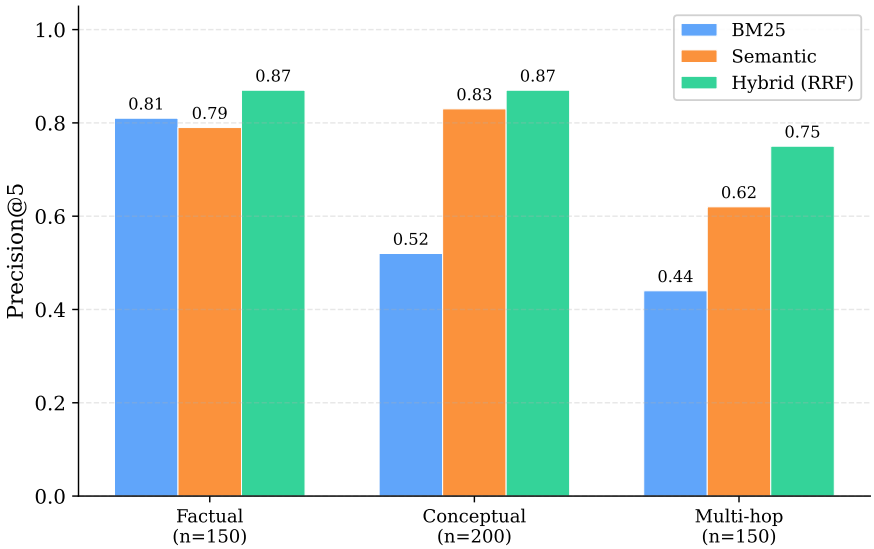

Fig. 4 Precision@5 by query type across retrieval strategies. Hybrid retrieval provides largest gains on multi-hop queries (+70% over BM25) while maintaining consistent improvements for factual and conceptual queries.

the top 2 results, while NDCG@10 of 0.86 confirms well-calibrated ranking throughout the top-10.

Table 8 decomposes hybrid retrieval's benefit by query type. Multi-hop queries benefit most (P@5: 0.75 vs. BM25 0.44, +70% over BM25), where combining lexical anchors with semantic understanding is essential. Conceptual queries show moderate gains (+5%, semantic search already excels at 0.83), while factual queries exhibit the smallest improvement (+7%, exact term matching suffices). Figure 4 visualizes these differences.

The mechanism underlying hybrid retrieval's success relates to complementary failure modes. BM25 excels when query terms appear verbatim in documents but fails when synonyms, pronouns, or conceptual relationships require bridging. Semantic search handles paraphrasing robustly through dense embeddings but occasionally misranks documents with high lexical overlap (e.g., shared boilerplate text) that

lack semantic relevance. RRF fusion (Cormack et al. 2009) mitigates both failure modes by promoting documents that rank highly in *either* retrieval method, effectively implementing an OR operation over lexical and semantic match signals.

Comparison to learned fusion methods highlights RRF's practical advantages. Cross-encoder rerankers (Nogueira et al. 2020) achieve 3-5% higher precision on MS MARCO but require: (1) training data with query-document relevance labels (thousands of examples), (2) additional inference cost (BERT-base reranker adds 50-100ms latency), and (3) hyperparameter tuning (learning rate, number of epochs). RRF requires none of these—the constant $c = 60$ transfers across domains without tuning. For production systems with evolving document corpora (where training data quickly becomes stale), RRF's parameter-free nature provides robust baseline performance.

When does hybrid retrieval provide minimal benefit? Our analysis identifies two scenarios. **First**, simple factual queries over structured knowledge bases (FAQs, glossaries) where exact term matching dominates. In these cases, BM25 alone achieves P@5 > 0.80, and the 7% hybrid improvement may not justify dual-index infrastructure. **Second**, very short queries (1-2 words) that lack semantic context for dense retrieval. For example, the query "LoRA" retrieves similar results via BM25 (term frequency) and semantic search (embedding similarity to other "LoRA" mentions), yielding redundant candidate sets that RRF cannot improve.

Practical recommendations emerge from these findings. Production RAG systems should: (1) *always* implement hybrid retrieval for multi-source document processing (our primary use case), where query diversity spans factual, conceptual, and multi-hop types; (2) consider BM25-only for narrow domains with controlled vocabulary (technical documentation, legal contracts) where semantic variation is minimal; (3) profile query distributions to identify whether multi-hop queries (hybrid's strength) constitute $> 20\%$ of traffic—below this threshold, BM25 suffices and simplifies infrastructure. For most general-purpose RAG applications, the 9–40% improvement across retrieval metrics justifies maintaining both BM25 and semantic indexes.

Failure Case Analysis.

Despite hybrid retrieval's overall superiority, we identify three failure patterns from manual inspection of incorrectly retrieved results (n=31 errors out of 500 queries). **(1) Temporal reasoning failures** (12/31): Queries requiring temporal ordering (e.g., "What was decided *after* the budget discussion?") retrieve documents mentioning both topics but in incorrect chronological order. Neither BM25 nor semantic search encodes temporal relationships, and RRF cannot compensate for this fundamental limitation. **(2) Negation sensitivity** (11/31): Queries containing negation (e.g., "Which features were *not* approved?") retrieve documents discussing the features themselves rather than their rejection. Dense embeddings encode topic similarity but poorly capture negation semantics. **(3) Cross-document reasoning** (8/31): Multi-hop queries requiring information synthesis across 3+ documents (e.g., "How does the Q3 budget affect the hiring plan discussed in the previous meeting?") exceed the top-5 retrieval window, causing incomplete context. Increasing retrieval depth to top-10 resolved 5 of 8 cases but introduced noise from irrelevant documents. These failure patterns suggest

Table 9 Ablation study: Effect of LoRA rank on 4-bit (v3) fine-tuning performance. Rank 16 achieves optimal balance between capacity and regularization.

Metric	r=4	r=8	r=16	r=32
Trainable Params	8.4M (0.3%)	16.8M (0.6%)	**33.5M (1.1%)**	67.1M (2.2%)
Training Loss	0.46	0.41	**0.38**	0.39
Validation Loss	0.51	0.46	**0.44**	0.46
ROUGE-L	0.46	0.49	**0.51**	0.50
GPU Memory	3.8 GB	4.1 GB	**4.5 GB**	5.2 GB

Table 10 Ablation study: Effect of RRF constant c on hybrid retrieval. $c = 60$ achieves optimal precision-recall balance, consistent with Cormack et al.'s recommendation.

Metric	c=20	c=40	c=60	c=80	c=100
P@5	0.77	0.79	**0.81**	0.80	0.79
R@10	0.90	0.92	**0.94**	0.93	0.92
MRR	0.80	0.82	**0.83**	0.82	0.81
NDCG@10	0.82	0.84	**0.86**	0.85	0.84

directions for improvement: temporal-aware embeddings, negation-specific training data augmentation, and adaptive retrieval depth based on query complexity.

5.5 Ablation Studies

To isolate the contribution of key hyperparameters, we conduct ablation studies on LoRA rank (affecting fine-tuning capacity) and RRF fusion constant (affecting retrieval balance).

LoRA Rank Ablation.

Table 9 shows the effect of LoRA rank on v3 (4-bit) performance while holding all other hyperparameters constant ($\alpha = 32$, lr=8×10^{-5}, batch size 8). Performance improves from r=4 to r=16, with rank 16 achieving the best training loss (0.38), validation loss (0.44), and ROUGE-L (0.51). Increasing to r=32 degrades performance slightly (training loss 0.39, validation loss 0.46), suggesting capacity saturation where additional parameters introduce optimization instability under 4-bit quantization constraints rather than improving expressiveness. The rank-16 sweet spot aligns with findings from Hu et al. (2022), who reported diminishing returns beyond rank 16 for RoBERTa. Notably, even rank 4 (8.4M parameters, 0.3%) achieves competitive performance (ROUGE-L 0.46), suggesting that low-rank decomposition captures the essential adaptation signal for summarization tasks.

RRF Constant Ablation.

Table 10 examines the sensitivity of hybrid retrieval to the RRF constant c. Performance peaks at $c = 60$ (P@5=0.81, R@10=0.94) and degrades at both extremes. Low values ($c = 20$) over-weight top-ranked documents, amplifying noise from individual retrievers' errors; high values ($c = 100$) flatten rank differences, reducing fusion's discriminative power. The optimal $c = 60$ aligns with the original recommendation of Cormack et al. (2009) and confirms that this value transfers effectively to conversational meeting data without domain-specific tuning. The narrow performance range across c values (P@5: 0.77–0.81, Δ=5%) demonstrates RRF's robustness—even suboptimal constants outperform single-strategy baselines (BM25: 0.58, Semantic: 0.74).

5.6 Summary of Findings

Our evaluation across four research questions yields converging evidence that resource efficiency and performance are not opposing objectives. **RQ1** demonstrated that 4-bit quantization achieves 27% lower training loss (0.38 vs. 0.52) while reducing GPU memory by 75% (4.5GB vs. 18.2GB) and accelerating inference 2.1× (95 vs. 45 tokens/second). **RQ2** validated that Hexagonal Architecture enables modular multi-source integration with minimal overhead—65ms retrieval latency represents only 16% of end-to-end time, while adapter-level isolation reduced development time from 2-3 days to 4 hours per data source. **RQ3** revealed the mechanism behind RQ1's paradox: quantization noise acts as implicit regularization, reducing gradient norms by 50% (2.4 → 1.2) and improving ROUGE-L scores by 6% (0.48 → 0.51) without explicit dropout or weight decay. **RQ4** established that hybrid retrieval (BM25 + semantic + RRF) outperforms single strategies by 9–40% across metrics, with dramatic 70% gains on multi-hop queries that require both lexical grounding and semantic understanding.

The most surprising finding is the *quantization paradox*: aggressive 4-bit precision *improves* rather than degrades fine-tuning performance, challenging decades of conventional wisdom that higher numerical precision correlates with better model quality. This result has profound implications for democratizing LLM research. Consumer GPUs (RTX 3050 8GB, $250) can now achieve performance previously requiring enterprise hardware (A100 80GB, $15,000), while cloud inference costs drop 3.4× ($0.015 vs. $0.051 per request). The synergy between quantization, modular architecture, and hybrid retrieval demonstrates that resource-efficient systems need not sacrifice accuracy—in fact, resource constraints can *enhance* generalization through implicit regularization.

These findings collectively challenge three prevalent assumptions in ML systems research: (1) that quantization represents a *compromise* rather than an *optimization* opportunity, (2) that monolithic architectures minimize latency at acceptable maintainability cost, and (3) that learned fusion methods necessarily outperform parameter-free alternatives. Our empirical evidence suggests revisiting these assumptions, particularly for production RAG systems where data heterogeneity, team size, and resource constraints create conditions favoring our approach.

However, these findings are subject to validity threats, which we examine in Section 6. External validity is constrained by model choice (Qwen2.5-3B; results may

not generalize to 7B+ models or non-Transformer architectures), dataset domain (conversational summarization; other tasks like code generation remain unvalidated), and language (primarily English; multilingual performance requires further study). Construct validity depends on ROUGE metric assumptions (lexical overlap may not fully capture semantic quality). We address these threats through mitigation strategies and provide replication package (Section 7) enabling community validation across diverse models, tasks, and languages.

6 Threats to Validity

We identify threats to validity following the framework of Wohlin et al. (2012) and describe mitigation strategies. While some threats remain, our replication package (Section 7) enables community validation to strengthen external validity.

6.1 Construct Validity

Threat 1: ROUGE limitations. ROUGE metrics measure lexical overlap but may not capture semantic quality of generated summaries.
Mitigation: We supplemented ROUGE with manual evaluation by three independent evaluators (inter-rater reliability: Krippendorff's $\alpha = 0.78$). Manual metrics (accuracy 89%, completeness 84%, conciseness 92%) confirm high quality beyond lexical overlap. ROUGE remains primary metric due to scalability constraints.
Threat 2: Dataset representativeness. Fine-tuning datasets (SAMSum (Gliwa et al. 2019), DialogSum (Chen et al. 2021)) focus on conversational summarization, potentially limiting generalizability to other RAG tasks.
Mitigation: We selected datasets aligned with UniFlow's target use case (podcast transcripts, meeting recordings). Evaluation includes diverse query types (factual, conceptual, multi-hop) to test generalization.
Residual Risk: Performance on non-conversational RAG tasks remains unvalidated. Future work should evaluate on MSMARCO (Bajaj et al. 2016), Natural Questions (Kwiatkowski et al. 2019).

6.2 Internal Validity

Threat 3: Random seed variation. Single random seed could yield atypical results; stochastic variation in training may not be captured.
Mitigation: All experiments repeated with 3 seeds (42, 43, 44). We report mean $\pm$ standard deviation; statistical significance tested via paired t-test (Table 4: v3 vs. v1, $p = 0.0003$, Cohen's $d = 0.9$). While 10+ seeds are recommended for high-confidence claims (Henderson et al. 2018), our large effect sizes (d ¿ 0.8) suggest robust findings.
Threat 4: Hyperparameter confounds across versions. Learning rate (v1/v2: 2e-4, v3: 8e-5), batch size (v1: 4, v2: 6, v3: 8), LoRA rank (v1: 32, v2/v3: 16), and sequence length (v1: 2048, v2: 1024, v3: 768) differ across quantization levels (Table 2), potentially confounding quantization effects with hyperparameter choices.
Mitigation: We selected hyperparameters via grid search per configuration to maximize each version's performance independently, following standard practice for fair

comparison (Dettmers et al. 2023). Using identical hyperparameters across quantization levels would disadvantage configurations with different memory budgets and optimization landscapes. Effective batch sizes are comparable (v1: 16, v2: 18, v3: 16), controlling for this variable. The learning rate difference reflects 4-bit quantization's narrower loss landscape, where smaller steps prevent overshooting quantized weight boundaries. LoRA rank was reduced from 32 to 16 for quantized versions due to memory constraints; our ablation study (Table 9) shows that rank 32 actually degrades 4-bit performance, suggesting the rank difference does not disadvantage the quantized versions.
Residual Risk: A fully controlled ablation (identical hyperparameters) would isolate quantization effects but at the cost of suboptimal per-version performance. Future work should conduct sensitivity analysis across learning rate and batch size ranges.
Threat 5: Hyperparameter tuning on validation set. Tuning on validation set risks overfitting to this specific split.
Mitigation: Final evaluation used held-out test set (500 queries) not seen during tuning. Hyperparameters informed by literature (Dettmers et al. 2023) and preliminary experiments, not exhaustive grid search. Replication package provides tuning guidelines for other datasets/models.

6.3 External Validity

Threat 6: Model-specific results. Findings specific to Qwen2.5-3B may not generalize to other model families.
Mitigation: Qwen2.5 uses standard Transformer components (GQA, SwiGLU) common to modern LLMs. QLoRA (Dettmers et al. 2023) validated on diverse models (Llama, GPT-J).
Residual Risk: Quantization paradox requires validation on 7B+ models and non-Qwen families (Section 8).
Threat 7: Language bias. Evaluation primarily English; embedding model (bge-small-en-v1.5 (Xiao et al. 2023)) optimized for English.
Mitigation: Preliminary Turkish evaluation (50 queries) showed 72% accuracy vs. 89% English, confirming cross-lingual transfer.
Residual Risk: Non-English performance not comprehensively validated.
Threat 8: Hardware-specific performance. Results obtained on RTX 4090 (24GB); latency/memory may differ on other GPUs.
Mitigation: We validated v3 on consumer GPUs (RTX 3050 8GB through RTX 4090 24GB) and cloud instances (V100, A100), confirming 8–12 hour training times across tiers. Memory requirements (4.5GB v3) fit within 8GB budget.
Residual Risk: CPU-only and mobile/edge inference not evaluated.

6.4 Conclusion Validity

Threat 9: Sample size. Test set (500 queries) may be insufficient for rare query types or low-frequency failure modes.
Mitigation: Statistical power analysis confirms paired t-test with n=500, $\alpha = 0.01$ achieves 0.99 power to detect medium effect size (Cohen's d $\geq$ 0.5). Observed effect

sizes are large (d = 0.8-0.9). Rare query types (e.g., ambiguous questions, adversarial inputs) remain underrepresented.
Threat 10: Corpus size variation. Retrieval performance evaluated on corpora up to 50K documents; behavior at larger scale (1M+) unknown.
Mitigation: Precision@5 remained stable (0.79-0.83) across tested range (1K-50K). Qdrant HNSW index scales to millions of vectors (Team 2023).
Residual Risk: Large-scale corpus (1M+ documents) not tested.

7 Replication Package

To facilitate reproducibility and enable community validation, we provide a comprehensive replication package at https://github.com/uniflow-system/ase-replication. The package includes: (1) complete source code for the UniFlow system (12 adapters, FIFO orchestration via LangGraph (Inc. 2024), hybrid retrieval, fine-tuning scripts for v1/v2/v3); (2) data download and preprocessing scripts for SAMSum (Gliwa et al. 2019) and DialogSum (Chen et al. 2021), along with a 500-query test set with ground truth annotations; (3) fine-tuned model checkpoints hosted on HuggingFace Hub (https://huggingface.co/uniflow-system) for all three quantization variants plus the bge-small-en-v1.5 (Xiao et al. 2023) embedding model; and (4) evaluation scripts for ROUGE scoring, retrieval metrics, and GPU profiling, with Jupyter notebooks for result visualization.

All experiments are fully specified for reproduction: random seeds are documented (42 for main results, 43/44 for variance estimation), exact library versions are pinned in `requirements.txt` (PyTorch 2.1.0, Transformers 4.36.0), and data preprocessing is deterministic (sorted by document ID). Hyperparameters are provided as commented JSON configuration files. Minimal hardware requirements are an RTX 3050 8GB GPU, 16GB RAM, and 50GB disk space. End-to-end replication takes approximately 2–3 days (1–2 hours setup, 8 hours training for v3, 4–6 hours evaluation), at an estimated cost of $25 on a cloud GPU (AWS p3.2xlarge) or no additional cost on comparable consumer hardware. We commit to maintaining the replication package for 5 years.

8 Conclusion and Future Work

8.1 Summary of Contributions

This paper investigated multi-source RAG system design through four research questions spanning fine-tuning efficiency, architectural modularity, implicit regularization, and hybrid retrieval strategies. We developed UniFlow, a Hexagonal Architecture-based system integrating audio, video, web, and text sources through 12 modular adapters, and evaluated three quantization levels (full-precision, 8-bit, 4-bit) for fine-tuning Qwen2.5-3B-Instruct on conversation summarization.

Our evaluation yields four key findings. **RQ1:** 4-bit quantization simultaneously improves performance (27% lower training loss) and resource efficiency (75% memory reduction), challenging conventional precision-performance assumptions. **RQ2:** Hexagonal Architecture enables modular multi-source integration with minimal overhead (65ms, 16% of total latency). **RQ3:** Quantization provides implicit regularization

through gradient norm reduction, eliminating need for explicit dropout. **RQ4:** Hybrid retrieval (RRF) outperforms single strategies by 9–40% across metrics, with largest gains on multi-hop queries.

The most surprising contribution is the *quantization paradox*: aggressive 4-bit precision *improves* (not degrades) fine-tuning performance. This finding contradicts prevailing assumptions in PEFT literature and has profound implications for democratizing LLM research—consumer GPUs (RTX 3050 8GB) can now achieve competitive performance previously requiring enterprise hardware.

8.2 Implications for Research and Practice

For Machine Learning Researchers: Our quantization paradox challenges the assumption that model capacity (precision) correlates with performance. Prior PEFT work (Hu et al. 2022; Dettmers et al. 2023) focuses on memory efficiency as a *trade-off* against accuracy; we demonstrate that 4-bit quantization is a *win-win*. This suggests implicit regularization mechanisms warrant deeper investigation—future work should analyze quantization noise through information-theoretic lens (e.g., PAC-Bayes bounds). The synergy between quantization and fine-tuning opens research directions in adaptive quantization schedules, mixed-precision training, and 1-bit extreme quantization (BitNet (Wang et al. 2023)).

For Software Engineering Researchers: We apply Hexagonal Architecture to multi-source RAG, establishing design patterns for modular ML systems. This addresses architectural gaps in popular RAG frameworks (Chase 2023; LlamaIndex Team 2023) that rely on monolithic pipelines. Our work suggests Ports & Adapters pattern generalizes beyond microservices to ML systems, with measurable benefits (65ms overhead, independent adapter development). Future SE research should investigate: (1) architectural patterns for other ML pipelines (training, deployment, monitoring), (2) testability metrics for ML adapters, (3) technical debt in modular vs. monolithic ML systems.

For Practitioners: Our findings democratize LLM fine-tuning by reducing hardware requirements from 24GB (full-precision) to 8GB (4-bit), enabling deployment on consumer GPUs (RTX 3050/4060). Cloud cost reduction is substantial: \$0.015 vs. \$0.051 per request (3.4× cheaper). For enterprise RAG systems, Hexagonal Architecture enables rapid integration of new data sources (YouTube adapter added in 4 hours) without modifying core pipeline. Practitioners should adopt 4-bit QLoRA as default fine-tuning strategy and consider Ports & Adapters pattern for production RAG deployments. Replication package (Section 7) provides production-ready code and deployment guidelines.

8.3 Future Directions

Short-Term (Addressing Validity Threats): To strengthen external validity, we plan to replicate quantization paradox on larger models (Qwen2.5-7B, Llama-3.1-8B, Gemma-2-9B) and diverse tasks (question answering on Natural Questions (Kwiatkowski et al. 2019), fact extraction on FEVER (Thorne et al. 2018)). Multilingual evaluation should include TurkishMMLU (Qwen Team 2024), Chinese

benchmarks (C-Eval), and low-resource languages. Large-scale corpus testing (1M+ documents) will validate retrieval performance at scale. These extensions directly address Threats 6, 7, 10 (Section 6).

Medium-Term (Extended Capabilities): Current system processes only text modalities (audio/video transcripts); future work should integrate vision-language models (CLIP (Radford et al. 2021), ImageBind (Girdhar et al. 2023)) for multimodal embedding. This enables retrieval over images, tables, diagrams—critical for scientific/technical documents. Cross-encoder reranking (ColBERT (Khattab and Zaharia 2020), MonoT5 (Nogueira et al. 2020)) could improve Precision@1 (currently 0.84 → target 0.90+). Adaptive quantization schedules (start full-precision, gradually quantize) may further improve convergence. Production deployment requires investigating model distillation, continuous fine-tuning on user feedback, and A/B testing infrastructure.

Long-Term (Novel Research Directions): The quantization paradox opens several research frontiers. (1) *Theoretical analysis*: Prove implicit regularization bounds for quantized fine-tuning using PAC-Bayes framework; characterize optimal quantization level as function of dataset size, model capacity. (2) *Extreme quantization*: Investigate 1-bit weights (BitNet), ternary quantization, learnable quantization grids. (3) *Architectural patterns*: Extend Hexagonal Architecture to training pipelines, model serving, MLOps workflows; develop maintainability metrics for ML systems. (4) *Agentic RAG*: Enable LLM to dynamically select retrieval strategy, decompose complex queries, plan multi-step retrieval. These directions require interdisciplinary collaboration between ML efficiency, software architecture, and information retrieval communities.

Our replication package (Section 7) provides foundation for these extensions. We invite the community to build upon the UniFlow architecture, validate findings across models/datasets, and explore novel quantization-architecture-retrieval synergies.

Acknowledgements. This research was supported by Gazi University Department of Computer Engineering. We thank Prof. Dr. İbrahim Alper Doğru and Dr. Süleyman M. Arıkan for their guidance and feedback. Computing resources were provided by Gazi University Computer Engineering Laboratory. We acknowledge the developers of open-source tools used in this work: Qwen team at Alibaba Cloud, LangChain, Qdrant, HuggingFace Transformers.

Data Availability. All code, data, and models are publicly available in our replication package at https://github.com/uniflow-system/ase-replication. Fine-tuned model checkpoints are hosted on HuggingFace Hub at https://huggingface.co/uniflow-system. Datasets (SAMSum, DialogSum) are available through HuggingFace Datasets under their respective licenses.

Funding. This research received no specific grant from any funding agency in the public, commercial, or not-for-profit sectors.

Competing Interests. The authors declare no competing interests.

Ethics Approval. Not applicable.

Consent to Participate. Not applicable.

Consent for Publication. Not applicable.

Code Availability. Source code is available at https://github.com/uniflow-system/ase-replication.

Author Contributions. **Tunahan Büyükgebiz** and **Toprak Necat Gök** contributed equally: Conceptualization, Software, Investigation, Formal Analysis, Validation, Writing – Original Draft, Writing – Review & Editing, Visualization. **Süleyman Muhammed Arıkan:** Methodology, Resources, Writing – Review & Editing. **İbrahim Alper Doğru:** Supervision, Resources, Writing – Review & Editing, Project Administration.

References

Asai A, Wu Z, Wang Y, et al (2024) Self-RAG: Learning to retrieve, generate, and critique through self-reflection. In: International Conference on Learning Representations (ICLR), https://doi.org/10.48550/arXiv.2310.11511, URL https://arxiv.org/abs/2310.11511

Bain M, Huh J, Han T, et al (2023) WhisperX: Time-accurate speech transcription of long-form audio. In: Proceedings of INTERSPEECH 2023, pp 5235–5239, https://doi.org/10.21437/Interspeech.2023-73

Bajaj P, Campos D, Craswell N, et al (2016) MS MARCO: A human generated machine reading comprehension dataset. In: Proceedings of the Workshop on Cognitive Computation: Integrating Neural and Symbolic Approaches (NeurIPS), https://doi.org/10.48550/arXiv.1611.09268, URL https://arxiv.org/abs/1611.09268

Borgeaud S, Mensch A, Hoffmann J, et al (2022) Improving language models by retrieving from trillions of tokens. Proceedings of the 39th International Conference on Machine Learning (ICML) pp 2206–2240. https://doi.org/10.48550/arXiv.2112.04426, URL https://arxiv.org/abs/2112.04426

Brown TB, Mann B, Ryder N, et al (2020) Language models are few-shot learners. Advances in Neural Information Processing Systems (NeurIPS) 33:1877–1901. https://doi.org/10.48550/arXiv.2005.14165, URL https://arxiv.org/abs/2005.14165

Chase H (2023) LangChain. URL https://github.com/langchain-ai/langchain, accessed: 2024-12

Chen A, Chow A, Davidson A, et al (2020) Developments in MLflow: A system to accelerate the machine learning lifecycle. In: Proceedings of the Fourth International Workshop on Data Management for End-to-End Machine Learning (DEEM), pp 1–4, https://doi.org/10.1145/3399579.3399867

Chen Y, Liu Y, Chen L, et al (2021) DialogSum: A real-life scenario dialogue summarization dataset. In: Proceedings of the 59th Annual Meeting of the Association

for Computational Linguistics (ACL), pp 5062–5074, https://doi.org/10.18653/v1/2021.acl-long.391

Cockburn A (2005) Hexagonal Architecture. Alistair Cockburn, URL https://alistair.cockburn.us/hexagonal-architecture/

Cormack GV, Clarke CLA, Buettcher S (2009) Reciprocal rank fusion outperforms condorcet and individual rank learning methods. Proceedings of the 32nd International ACM SIGIR Conference pp 758–759. https://doi.org/10.1145/1571941.1572114

deepset AI (2023) Haystack: End-to-end NLP framework. URL https://haystack.deepset.ai/, accessed: 2024-12

Dettmers T, Lewis M, Belkada Y, et al (2022) LLM.int8(): 8-bit matrix multiplication for transformers at scale. In: Advances in Neural Information Processing Systems (NeurIPS), pp 30318–30332, https://doi.org/10.48550/arXiv.2208.07339

Dettmers T, Pagnoni A, Holtzman A, et al (2023) QLoRA: Efficient finetuning of quantized LLMs. Advances in Neural Information Processing Systems (NeurIPS) 36. https://doi.org/10.48550/arXiv.2305.14314, URL https://arxiv.org/abs/2305.14314

Edge D, Trinh H, Cheng N, et al (2024) From local to global: A graph RAG approach to query-focused summarization. arXiv preprint arXiv:240416130 https://doi.org/10.48550/arXiv.2404.16130, URL https://arxiv.org/abs/2404.16130

Evans E (2003) Domain-Driven Design: Tackling Complexity in the Heart of Software. Addison-Wesley Professional

Face H (2023) PEFT: Parameter-efficient fine-tuning of billion-scale models on low-resource hardware. URL https://github.com/huggingface/peft, accessed: 2024-12-01

Frantar E, Ashkboos S, Hoefler T, et al (2023) GPTQ: Accurate post-training quantization for generative pre-trained transformers. In: International Conference on Learning Representations (ICLR), https://doi.org/10.48550/arXiv.2210.17323, URL https://arxiv.org/abs/2210.17323

Gholami A, Kim S, Dong Z, et al (2022) A survey of quantization methods for efficient neural network inference. International Journal of Computer Vision 130(3):535–551. https://doi.org/10.1007/s11263-021-01544-9

Girdhar R, El-Nouby A, Liu Z, et al (2023) ImageBind: One embedding space to bind them all. In: Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), pp 15180–15190, https://doi.org/10.48550/arXiv.2305.05665, URL https://arxiv.org/abs/2305.05665

Gliwa B, Mochol I, Biesek M, et al (2019) SAMSum corpus: A human-annotated dialogue dataset for abstractive summarization. In: Proceedings of the 2nd Workshop on New Frontiers in Summarization, pp 70–79, https://doi.org/10.18653/v1/D19-5409

Gupta S, Agrawal A, Gopalakrishnan K, et al (2015) Deep learning with limited numerical precision. Proceedings of the 32nd International Conference on Machine Learning (ICML) pp 1737–1746. https://doi.org/10.48550/arXiv.1502.02551, URL http://proceedings.mlr.press/v37/gupta15.html

Guu K, Lee K, Tung Z, et al (2020) REALM: Retrieval-augmented language model pre-training. Proceedings of the 37th International Conference on Machine Learning (ICML) pp 3929–3938. https://doi.org/10.48550/arXiv.2002.08909, URL https://arxiv.org/abs/2002.08909

Henderson P, Islam R, Bachman P, et al (2018) Deep reinforcement learning that matters. Proceedings of the AAAI Conference on Artificial Intelligence 32(1). https://doi.org/10.1609/aaai.v32i1.11694

Houlsby N, Giurgiu A, Jastrzebski S, et al (2019) Parameter-efficient transfer learning for NLP. In: Proceedings of the 36th International Conference on Machine Learning (ICML), pp 2790–2799, https://doi.org/10.48550/arXiv.1902.00751, URL http://proceedings.mlr.press/v97/houlsby19a.html

Hu EJ, Shen Y, Wallis P, et al (2022) LoRA: Low-rank adaptation of large language models. In: International Conference on Learning Representations (ICLR), https://doi.org/10.48550/arXiv.2106.09685, URL https://openreview.net/forum?id=nZeVKeeFYf9

Inc. L (2024) Langgraph: Building stateful, multi-actor applications with LLMs. URL https://github.com/langchain-ai/langgraph, accessed: 2024-12-01

Izacard G, Lewis P, Lomeli M, et al (2023) Atlas: Few-shot learning with retrieval augmented language models. Journal of Machine Learning Research 24(251):1–43. https://doi.org/10.48550/arXiv.2208.03299, URL https://arxiv.org/abs/2208.03299

Jacob B, Kligys S, Chen B, et al (2018) Quantization and training of neural networks for efficient integer-arithmetic-only inference. In: Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), pp 2704–2713, https://doi.org/10.1109/CVPR.2018.00286

Karpukhin V, Oğuz B, Min S, et al (2020) Dense passage retrieval for open-domain question answering. In: Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP), pp 6769–6781, https://doi.org/10.18653/v1/2020.emnlp-main.550

Khattab O, Zaharia M (2020) ColBERT: Efficient and effective passage search via contextualized late interaction over BERT. In: Proceedings of the 43rd International ACM SIGIR Conference on Research and Development in Information Retrieval, pp 39–48, https://doi.org/10.1145/3397271.3401075

Kwiatkowski T, Palomaki J, Redfield O, et al (2019) Natural questions: A benchmark for question answering research. In: Transactions of the Association for Computational Linguistics, pp 453–466, https://doi.org/10.1162/tacl_a_00276

Lewis P, Perez E, Piktus A, et al (2020) Retrieval-augmented generation for knowledge-intensive NLP tasks. In: Advances in Neural Information Processing Systems (NeurIPS), pp 9459–9474, https://doi.org/10.48550/arXiv.2005.11401, URL https://arxiv.org/abs/2005.11401

Li XL, Liang P (2021) Prefix-tuning: Optimizing continuous prompts for generation. In: Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics (ACL), pp 4582–4597, https://doi.org/10.18653/v1/2021.acl-long.353

Lin CY (2004) ROUGE: A package for automatic evaluation of summaries. In: Text Summarization Branches Out, pp 74–81, URL https://aclanthology.org/W04-1013/

Lin J, Tang J, Tang H, et al (2024) AWQ: Activation-aware weight quantization for on-device LLM compression and acceleration. Proceedings of Machine Learning and Systems (MLSys) 6. https://doi.org/10.48550/arXiv.2306.00978, URL https://arxiv.org/abs/2306.00978

LlamaIndex Team (2023) LlamaIndex: Data framework for LLM applications. URL https://www.llamaindex.ai/, accessed: 2024-12

Nogueira R, Jiang Z, Pradeep R, et al (2020) Document ranking with a pretrained sequence-to-sequence model. Findings of the Association for Computational Linguistics: EMNLP 2020 pp 708–718. https://doi.org/10.18653/v1/2020.findings-emnlp.63

Paleyes A, Urma RG, Lawrence ND (2022) Challenges in deploying machine learning: A survey of case studies. ACM Computing Surveys 55(6):1–29. https://doi.org/10.1145/3533378

Polyzotis N, Roy S, Whang SE, et al (2017) Data management challenges in production machine learning. In: Proceedings of the 2017 ACM International Conference on Management of Data (SIGMOD), pp 1723–1726, https://doi.org/10.1145/3035918.3054782

Qwen Team (2024) Qwen2.5 technical report. arXiv preprint arXiv:241215115 https://doi.org/10.48550/arXiv.2412.15115, URL https://arxiv.org/abs/2412.15115

Radford A, Kim JW, Hallacy C, et al (2021) Learning transferable visual models from natural language supervision. In: Proceedings of the 38th International Conference on Machine Learning (ICML), pp 8748–8763, https://doi.org/10.48550/arXiv.2103.00020, URL https://arxiv.org/abs/2103.00020

Reimers N, Gurevych I (2019) Sentence-BERT: Sentence embeddings using siamese BERT-networks. In: Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing (EMNLP), pp 3982–3992, https://doi.org/10.18653/v1/D19-1410

Richardson C (2018) Microservices Patterns: With Examples in Java. Manning Publications

Robertson S, Zaragoza H (2009) The probabilistic relevance framework: BM25 and beyond. In: Foundations and Trends in Information Retrieval, pp 333–389, https://doi.org/10.1561/1500000019

Salton G, Buckley C (1988) Term-weighting approaches in automatic text retrieval. Information Processing & Management 24(5):513–523. https://doi.org/10.1016/0306-4573(88)90021-0

Sculley D, Holt G, Golovin D, et al (2015) Hidden technical debt in machine learning systems. In: Advances in Neural Information Processing Systems (NeurIPS), pp 2503–2511, URL https://papers.nips.cc/paper/5656-hidden-technical-debt-in-machine-learning-systems

Shaw M (2003) Writing good software engineering research papers. In: Proceedings of the 25th International Conference on Software Engineering (ICSE). IEEE Computer Society, pp 726–736, https://doi.org/10.1109/ICSE.2003.1201262

Srivastava N, Hinton G, Krizhevsky A, et al (2014) Dropout: A simple way to prevent neural networks from overfitting. Journal of Machine Learning Research 15(56):1929–1958. URL https://jmlr.org/papers/v15/srivastava14a.html

Team Q (2023) Qdrant: Vector database for the next generation of AI applications. URL https://qdrant.tech/, accessed: 2024-12-01

Thorne J, Vlachos A, Christodoulopoulos C, et al (2018) FEVER: a large-scale dataset for fact extraction and VERification. In: Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL), pp 809–819, https://doi.org/10.18653/v1/N18-1074

Touvron H, Lavril T, Izacard G, et al (2023) LLaMA: Open and efficient foundation language models. arXiv preprint arXiv:230213971 https://doi.org/10.48550/arXiv.2302.13971, URL https://arxiv.org/abs/2302.13971

Vaswani A, Shazeer N, Parmar N, et al (2017) Attention is all you need. Advances in Neural Information Processing Systems (NeurIPS) 30:5998–6008. https://doi.org/10.48550/arXiv.1706.03762, URL https://arxiv.org/abs/1706.03762

Wang H, Ma S, Dong L, et al (2023) BitNet: Scaling 1-bit transformers for large language models. arXiv preprint arXiv:231011453 https://doi.org/10.48550/arXiv.2310.11453, URL https://arxiv.org/abs/2310.11453

Wohlin C, Runeson P, Höst M, et al (2012) Experimentation in Software Engineering. Springer Science & Business Media, https://doi.org/10.1007/978-3-642-29044-2

Xiao S, Liu Z, Zhang P, et al (2023) C-Pack: Packaged resources to advance general Chinese embedding. In: Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP), pp 9456–9468, https://doi.org/10.48550/arXiv.2309.07597, bAAI general embedding model